



JNAN VIKAS MANDAL'S
PADMASHREE DR. R.T.DOSHI DEGREE COLLEGE OF
INFORMATION TECHNOLOGY

MOHANLAL RAICHAND MEHTA COLLEGE OF
COMMERCE

DIWALIMAA DEGREE COLLEGE OF SCIENCE

CERTIFICATE

This is to certify that the Mr./Miss. _____
having Roll No _____ of B.Sc. CS Semester V has completed
the practical work in the subject of **Artificial Intelligence** under the guidance of
Prof Janhavi Kshirsagar during the Academic year 2023-24 being the partial
requirement for the fulfillment of the curriculum of Degree of Bachelor of Science
in Computer Science, University of Mumbai.

Place:

Date:

Sign of Subject In Charge

Sign of External Examiner

Sign CS-IT Coordinator

T.Y.B.Sc. COMPUTER SCIENCE
SEMESTER: V

Artificial Intelligence
(USCSP501)

Journal

Index

Sr.no	Practical	Date	Sign
1.	<ul style="list-style-type: none">• Breadth First Search & Iterative Depth First Search<ul style="list-style-type: none">○ Implement the Breadth First Search algorithm to solve a given problem.○ Implement the Iterative Depth First Search algorithm to solve the same problem.• Compare the performance and efficiency of both algorithms.		
2.	<ul style="list-style-type: none">• A* Search and Recursive Best-First Search<ul style="list-style-type: none">○ Implement A* Search algorithm to solve a given problem.○ Implement Recursive Best-First Search algorithm to solve the same problem.• Compare the performance and effectiveness of both algorithms.		
3.	<ul style="list-style-type: none">• Decision Tree Learning<ul style="list-style-type: none">○ Implement the Decision Tree Learning algorithm to build a decision tree for a given dataset.○ Evaluate the accuracy and effectiveness of the decision tree on test data.• Visualize and interpret the generated decision tree.		
4.	<ul style="list-style-type: none">• Feed Forward Backpropagation Neural Network<ul style="list-style-type: none">○ Implement the Feed Forward Backpropagation algorithm to train a neural network.○ Use a given dataset to train the neural network for a specific task.• Evaluate the performance of the trained network on test data.		
5.	<ul style="list-style-type: none">• Support Vector Machines (SVM)<ul style="list-style-type: none">○ Implement the SVM algorithm for binary classification.○ Train an SVM model using a given dataset and optimize its parameters.• Evaluate the performance of the SVM model on		

	test data and analyze the results.		
6.	<ul style="list-style-type: none"> • Adaboost Ensemble Learning <ul style="list-style-type: none"> ○ Implement the Adaboost algorithm to create an ensemble of weak classifiers. ○ Train the ensemble model on a given dataset and evaluate its performance. ○ Compare the results with individual weak classifiers. 		
7.	<ul style="list-style-type: none"> • Naive Bayes' Classifier <ul style="list-style-type: none"> ○ Implement the Naive Bayes' algorithm for classification. ○ Train a Naive Bayes' model using a given dataset and calculate class probabilities. • Evaluate the accuracy of the model on test data and analyze 		
8.	<ul style="list-style-type: none"> • K-Nearest Neighbors (K-NN) <ul style="list-style-type: none"> ○ Implement the K-NN algorithm for classification or regression. ○ Apply the K-NN algorithm to a given dataset and predict the class or value for test data. • Evaluate the accuracy or error of the predictions and analyze the results. 		


PRACTICAL 1

A. Aim: Implement Breadth First Search Algorithm

Dataset: [RMP.py File](#)

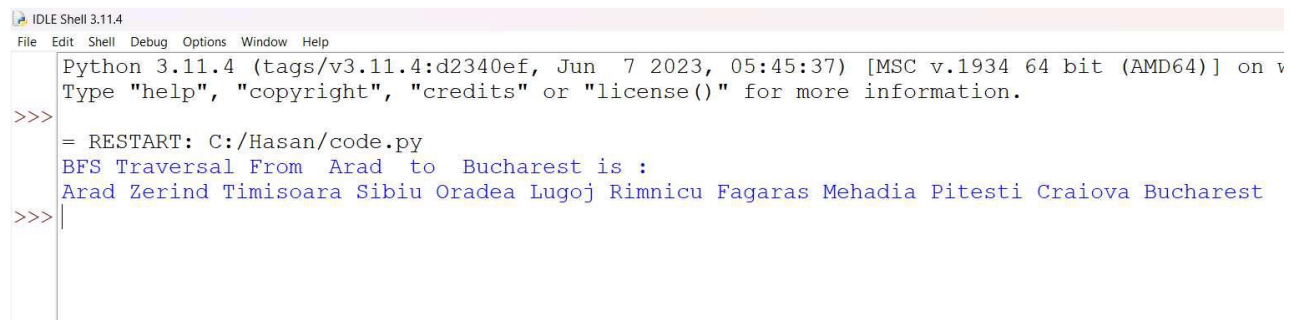
Requirement: RMP.py, Python IDLE

Diagram:



A. Code: Implement the Breadth First Search algorithm to solve a given problem.

```
import queue as Q
from RMP import dict_gn
start = 'Arad'
goal = "Bucharest"
result=""
def BFS(city,cityq,visitedq):
    global result
    if city==start:
        result = result + "" + city
    for eachcity in dict_gn[city].keys():
        if eachcity==goal:
            result = result + " " + eachcity
            return
        if eachcity not in cityq.queue and eachcity not in visitedq.queue:
            cityq.put(eachcity)
            result = result + " " + eachcity
    visitedq.put(city)
    BFS(cityq.get(),cityq,visitedq)
def main():
    cityq = Q.Queue()
    visitedq = Q.Queue()
    BFS(start,cityq,visitedq)
    print("BFS Traversal From ", start," to " , goal, "is :")
    print(result)
main()
```



```
IDLE Shell 3.11.4
File Edit Shell Debug Options Window Help
Python 3.11.4 (tags/v3.11.4:d2340ef, Jun 7 2023, 05:45:37) [MSC v.1934 64 bit (AMD64)] on v
Type "help", "copyright", "credits" or "license()" for more information.
>>>
= RESTART: C:/Hasan/code.py
BFS Traversal From Arad to Bucharest is :
Arad Zerind Timisoara Sibiu Oradea Lugoj Rimnicu Fagaras Mehadia Pitesti Craiova Bucharest
>>>
```

B. Code: Implement the Iterative Depth First Search algorithm to solve the same problem.

```
import queue as Q
from RMP import dict_gn
start = "Arad"
goal = "Bucharest"
result = ""

def DLS(city,visitedstack,startlimit,endlimit):
    global result
    found = 0
    result = result + city + " "
    visitedstack.append(city)
    if city == goal:
        return 1
    if startlimit == endlimit:
        return 0
    for eachcity in dict_gn[city].keys():
        if eachcity not in visitedstack:
            found = DLS(eachcity,visitedstack,startlimit+1,endlimit)
    if found:
        return found

def IDDFS(city,visitedstack,endlimit):
    global result
    for i in range(0,endlimit):
        print("Seaching at Limit:", i)
        found = DLS(city,visitedstack, 0 , i)
        if found:
            print("Found")
            break
    else:
        print("Not Found!")
        print(result)
        print("_____")
        result=""
        visitedstack = []
```

```
def main():
    visitedstack = []
    IDDFS(start,visitedstack,9)
    print("IDDFS Traversal from ", start, " to ",goal," is:")
    print(result)
main()
```

Output:

```
>>> Type "help", "copyright", "credits" or "license()" for more information.
>>> = RESTART: C:\Hasan\code.py
Seaching at Limit: 0
Not Found!
Arad

Seaching at Limit: 1
Not Found!
Arad Zerind Timisoara Sibiu

Seaching at Limit: 2
Not Found!
Arad Zerind Oradea Timisoara Lugoj Sibiu Rimnicu Fagaras

Seaching at Limit: 3
Not Found!
Arad Zerind Oradea Sibiu Timisoara Lugoj Mehadia

Seaching at Limit: 4
Not Found!
Arad Zerind Oradea Sibiu Rimnicu Fagaras Timisoara Lugoj Mehadia Drobeta

Seaching at Limit: 5
Found
IDDFS Traversal from Arad to Bucharest is:
Arad Zerind Oradea Sibiu Rimnicu Pitesti Craiova Fagaras Bucharest
>>>
```


PRACTICAL 2

AIM: A* Search and Recursive Best-First Search

Dataset: [RMP.py File](#)

Code: Implement the A* Search algorithm for solving a pathfinding problem.

```
import queue as Q
from RMP import dict_gn
from RMP import dict_hn

start = 'Arad'
goal = 'Bucharest'
result = ""

def get_fn(citystr):
    cities=citystr.split(",")
    hn=gn=0
    for ctr in range(0, len(cities)-1):
        gn=gn+dict_gn[cities[ctr]][cities[ctr+1]]
    hn=dict_hn[cities[len(cities)-1]]
    return(hn+gn)
def expand(cityq):
    global result
    tot, citystr, thiscity=cityq.get()
    if thiscity==goal:
        result=citystr+"::"+str(tot)
        return
    for cty in dict_gn[thiscity]:
        cityq.put((get_fn(citystr+","+cty),citystr+","+cty,cty))
    expand(cityq)
def main():
    cityq=Q.PriorityQueue()
    thiscity=start
    cityq.put((get_fn(start),start,thiscity))
    expand(cityq)
    print("The A* path with the total is: ")
    print(result)
main()
```

Output:

```
IDLE Shell 3.11.4
File Edit Shell Debug Options Window Help
Python 3.11.4 (tags/v3.11.4:d2340ef, Jun 7 202
Type "help", "copyright", "credits" or "license
>>>
= RESTART: C:/Hasan/code.py
The A* path with the total is:
Arad,Sibiu,Rimnicu,Pitesti,Bucharest::418
>>> |
```

Code: Implement the Recursive Best-First Search algorithm for the same problem.

```
import queue as Q
from RMP import dict_gn
from RMP import dict_hn

start = 'Arad'
goal = 'Bucharest'
result = ""

def get_fn(citystr):
    cities=citystr.split(",")
    hn=gn=0
    for ctr in range(0, len(cities)-1):
        gn=gn+dict_gn[cities[ctr]][cities[ctr+1]]
    hn=dict_hn[cities[len(cities)-1]]
    return(hn+gn)

def printout(cityq):
    for i in range(0, cityq.qsize()):
        print(cityq.queue[i])

def expand(cityq):
    global result
    tot, citystr, thiscity = cityq.get()
    nexttot = 999
    if not cityq.empty():
        nexttot,nextcitystr,nextthiscity=
        cityq.queue[0]
    if thiscity== goal and tot < nexttot:
        result = citystr + "::" + str(tot)
        return
    print("Expanded city ----- ", thiscity)
    print("second best f(n) ----- ", nexttot)
    tempq = Q.PriorityQueue()
    for cty in dict_gn[thiscity]:
        tempq.put((get_fn(citystr+' '+cty), citystr+' '+cty, cty))
        for ctr in range(1,3):
            ctrtot, ctrcitystr, ctrthiscity = tempq.get()
            if ctrtot < nexttot:
                cityq.put((ctrtot, ctrcitystr, ctrthiscity))
            else:
                cityq.put((ctrtot, citystr, thiscity))
        break
    printout(cityq)
    expand(cityq)
```

```

def main():
    cityq=Q.PriorityQueue()
    thiscity=start
    cityq.put((999, "NA", "NA"))
    cityq.put((get_fn(start), start, thiscity))
    expand(cityq)
    print(result)
main()

```

Output:

```

>>> = RESTART: C:\Hasan\code.py
Expaded city ----- Arad
second best f(n)----- 999
(393, 'Arad,Sibiu', 'Sibiu')
(999, 'NA', 'NA')
(447, 'Arad,Timisoara', 'Timisoara')
Expaded city ----- Sibiu
second best f(n)----- 447
(413, 'Arad,Sibiu,Rimnicu', 'Rimnicu')
(415, 'Arad,Sibiu,Fagaras', 'Fagaras')
(447, 'Arad,Timisoara', 'Timisoara')
(999, 'NA', 'NA')
Expaded city ----- Rimnicu
second best f(n)----- 415
(415, 'Arad,Sibiu,Fagaras', 'Fagaras')
(417, 'Arad,Sibiu,Rimnicu', 'Rimnicu')
(447, 'Arad,Timisoara', 'Timisoara')
(999, 'NA', 'NA')
Expaded city ----- Fagaras
second best f(n)----- 417
(417, 'Arad,Sibiu,Rimnicu', 'Rimnicu')
(450, 'Arad,Sibiu,Fagaras', 'Fagaras')
(447, 'Arad,Timisoara', 'Timisoara')
(999, 'NA', 'NA')
Expaded city ----- Rimnicu
second best f(n)----- 447
(417, 'Arad,Sibiu,Rimnicu,Pitesti', 'Pitesti')
(447, 'Arad,Timisoara', 'Timisoara')
(999, 'NA', 'NA')
(450, 'Arad,Sibiu,Fagaras', 'Fagaras')
(526, 'Arad,Sibiu,Rimnicu', 'Rimnicu')
Expaded city ----- Pitesti
second best f(n)----- 447
(418, 'Arad,Sibiu,Rimnicu,Pitesti,Bucharest', 'Bucharest')
(447, 'Arad,Timisoara', 'Timisoara')
(607, 'Arad,Sibiu,Rimnicu,Pitesti', 'Pitesti')
(526, 'Arad,Sibiu,Rimnicu', 'Rimnicu')
(450, 'Arad,Sibiu,Fagaras', 'Fagaras')
(999, 'NA', 'NA')
Arad,Sibiu,Rimnicu,Pitesti,Bucharest::418
>>>

```

PRACTICAL NO: 3

Aim: Implement the decision tree learning algorithm to build a decision tree for a given dataset. Evaluate the accuracy and efficiency on the test data set.

Implementing Decision Tree using Scikit Learn

This notebook is a reference notebook to a blog, [Decision Tree for Beginners](#).

```
In [1]: #numpy and pandas initialization
import numpy as np
import pandas as pd
```

```
In [2]: #Loading the PlayTennis data
PlayTennis = pd.read_csv("../input/PlayTennis.csv")
```

```
In [3]: PlayTennis
```

```
Out[3]:
```

	outlook	temp	humidity	windy	play
0	sunny	hot	high	False	no
1	sunny	hot	high	True	no
2	overcast	hot	high	False	yes
3	rainy	mild	high	False	yes
4	rainy	cool	normal	False	yes
5	rainy	cool	normal	True	no
6	overcast	cool	normal	True	yes
7	sunny	mild	high	False	no
8	sunny	cool	normal	False	yes
9	rainy	mild	normal	False	yes
10	sunny	mild	normal	True	yes
11	overcast	mild	high	True	yes
12	overcast	hot	normal	False	yes
13	rainy	mild	high	True	no

It is easy to implement Decision Tree with numerical values. We can convert all the non numerical values into numerical values using [LabelEncoder](#)

```
In [4]: from sklearn.preprocessing import LabelEncoder
Le = LabelEncoder()

PlayTennis['outlook'] = Le.fit_transform(PlayTennis['outlook'])
PlayTennis['temp'] = Le.fit_transform(PlayTennis['temp'])
PlayTennis['humidity'] = Le.fit_transform(PlayTennis['humidity'])
PlayTennis['windy'] = Le.fit_transform(PlayTennis['windy'])
PlayTennis['play'] = Le.fit_transform(PlayTennis['play'])
```

```
In [5]: PlayTennis
```

```
Out[5]:
```

	outlook	temp	humidity	windy	play
0	2	1	0	0	0
1	2	1	0	1	0
2	0	1	0	0	1
3	1	2	0	0	1
4	1	0	1	0	1
5	1	0	1	1	0
6	0	0	1	1	1
7	2	2	0	0	0
8	2	0	1	0	1
9	1	2	1	0	1
10	2	2	1	1	1
11	0	2	0	1	1
12	0	1	1	0	1
13	1	2	0	1	0

- Lets split the training data and its corresponding prediction values.
- y - holds all the decisions.
- X - holds the training data.

```
In [6]: y = PlayTennis['play']  
X = PlayTennis.drop(['play'],axis=1)
```

```
In [7]: # Fitting the model  
from sklearn import tree  
clf = tree.DecisionTreeClassifier(criterion = 'entropy')  
clf = clf.fit(X, y)
```

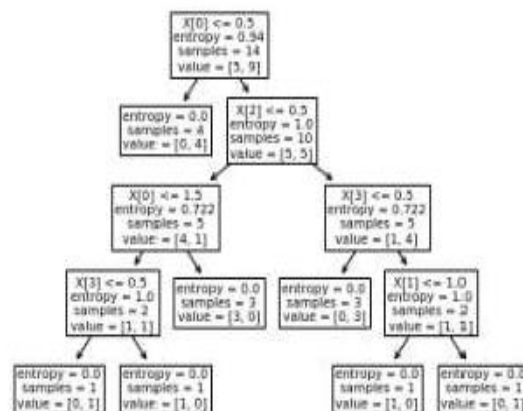
```
In [8]: # We can visualize the tree using tree.plot_tree  
tree.plot_tree(clf)
```



```

Out[8]: [Text(133.92000000000002, 195.696, 'X[0] <= 0.5\nentropy = 0.94\nsamples = 14\nvalue = [5, 9]'),
Text(100.44000000000001, 152.208, 'entropy = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(167.40000000000003, 152.208, 'X[2] <= 0.5\nentropy = 1.0\nsamples = 10\nvalue = [5, 5]'),
Text(100.44000000000001, 108.72, 'X[0] <= 1.5\nentropy = 0.722\nsamples = 5\nvalue = [4, 1]'),
Text(66.96000000000001, 65.232, 'X[3] <= 0.5\nentropy = 1.0\nsamples = 2\nvalue = [1, 1]'),
Text(33.480000000000004, 21.744, 'entropy = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(100.44000000000001, 21.744, 'entropy = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(133.92000000000002, 65.232, 'entropy = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(234.36, 108.72, 'X[3] <= 0.5\nentropy = 0.722\nsamples = 5\nvalue = [1, 4]'),
Text(200.88000000000002, 65.232, 'entropy = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(267.84000000000003, 65.232, 'X[1] <= 1.0\nentropy = 1.0\nsamples = 2\nvalue = [1, 1]'),
Text(234.36, 21.744, 'entropy = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(301.32000000000005, 21.744, 'entropy = 0.0\nsamples = 1\nvalue = [0, 1]')]

```



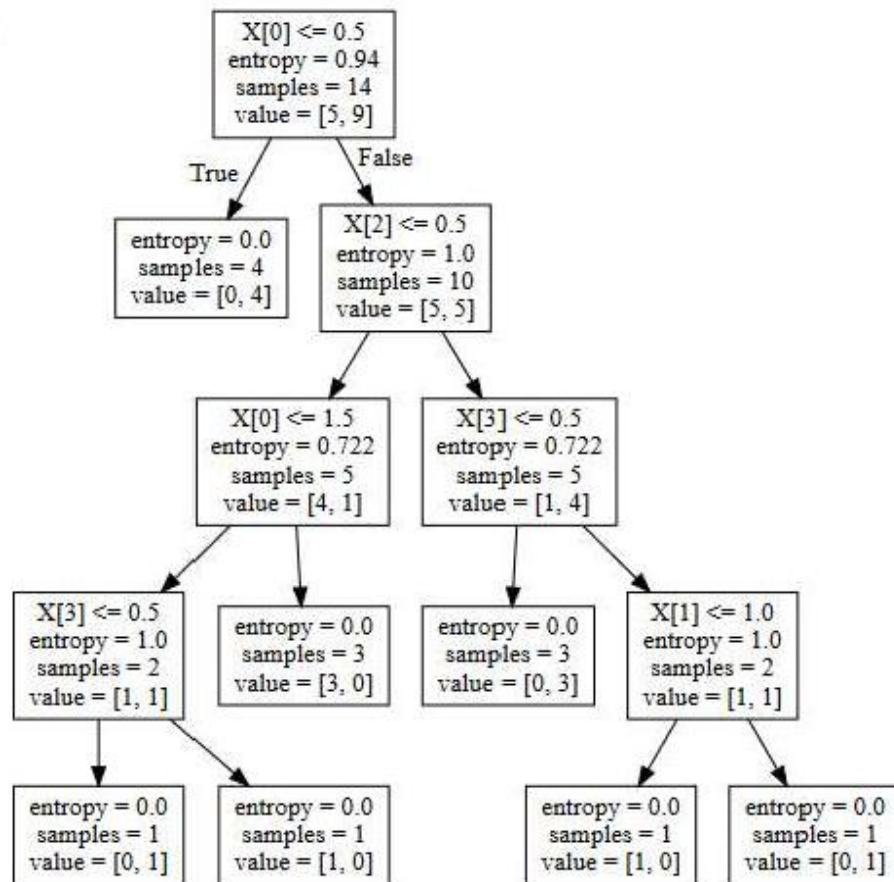
GraphViz gives a better and clearer Graph.

```

In [9]: import graphviz
dot_data = tree.export_graphviz(clf, out_file=None)
graph = graphviz.Source(dot_data)
graph


```

Out[9]:



In the above graph,

- X[0] -> Outlook
- X[1] -> Temperature
- X[2] -> Humidity
- X[3] -> Wind



values

Since we don't have any data to test, we can just make the model to predict our train data.

```
In [10]: # The predictions are stored in X_pred
X_pred = clf.predict(X)
```



```
In [11]: # verifying if the model has predicted it all right.  
x_pred == y
```

```
Out[11]: 0    True  
        1    True  
        2    True  
        3    True  
        4    True  
        5    True  
        6    True  
        7    True  
        8    True  
        9    True  
       10    True  
       11    True  
       12    True  
       13    True  
Name: play, dtype: bool
```

PRACTICAL NO: 4

AIM: Feed Forward Back propagation Neural Network

- Implement the Feed Forward Back propagation algorithm to train a neural network
- Use a given dataset to train the neural network for a specific task

Requirement: Python IDLE

Code:

```
from doctest import OutputChecker
import numpy as np
class NeuralNetwork():
    def __init__(self):
        np.random.seed()
        self.synaptic_weights=2*np.random.random((3,1))-1
    def sigmoid(self,x):
        return 1/(1+np.exp(-x))
    def sigmoid_derivative(self,x):
        return x*(1-x)
    def train(self,training_inputs,training_outputs,training_iterations):
        for iteration in range(training_iterations):
            output=self.think(training_inputs)
            error = training_outputs-output
            adjustments=np.dot(training_inputs.T,error*self.sigmoid_derivative(output))
            self.synaptic_weights +=adjustments
    def think(self,inputs):
        inputs=inputs.astype(float)
        output=self.sigmoid(np.dot(inputs,self.synaptic_weights))
        return output

if __name__ == "__main__":
    #initializing the neuron class
    neural_network = NeuralNetwork()
    print("Beginning Randomly Generated Weights: ")
    print(neural_network.synaptic_weights)
    #training data consisting of 4 examples--3 input values and 1 output
    training_inputs = np.array([[0,0,1],
    [1,1,1],
    [1,0,1],
    [0,1,1]])
```

```
training_outputs = np.array([[0,1,1,0]]).T
#training taking place
neural_network.train(training_inputs, training_outputs, 15000)
print("Ending Weights After Training: ")
print(neural_network.synaptic_weights)
user_input_one = str(input("User Input One: "))
user_input_two = str(input("User Input Two: "))
user_input_three = str(input("User Input Three: "))
print("Considering New Situation: ", user_input_one, user_input_two, user_input_three)
print("New Output data: ")
print(neural_network.think(np.array([user_input_one, user_input_two, user_input_three])))
```

Output:

```
Beginning Randomly Generated Weights:
[[0.18138631]
 [0.03957296]
 [0.68171289]]
Ending Weights After Training:
[[10.08723627]
 [-0.20745403]
 [-4.83719347]]
User Input One: 2
User Input Two: 3
User Input Three: 2
Considering New Situation: 2 3 2
New Output data:
[0.9999487]
```

PRACTICAL NO: 5

Aim: Implement the SVM algorithm for binary classification. Train a SVM Model using the given dataset. Evaluate the performance on test data and analyze the results.

Importing Libraries

```
In [1]: from warnings import filterwarnings
        filterwarnings("ignore")
```

```
In [2]: pip install skompiler
```

```
Collecting skompiler
  Downloading SKompiler-0.6.tar.gz (45 kB)
    |#####| 45 kB 388 kB/s
Requirement already satisfied: scikit-learn>=0.22 in /opt/conda/lib/python3.7/site-packages (from skompiler) (0.23.2)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.22->skompiler) (1.0.1)
Requirement already satisfied: scipy>=0.19.1 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.22->skompiler) (1.6.3)
Requirement already satisfied: numpy>=1.13.3 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.22->skompiler) (1.19.5)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.22->skompiler) (2.1.0)
Building wheels for collected packages: skompiler
  Building wheel for skompiler (setup.py) ... -E E\B Edone
  Created wheel for skompiler: filename=SKompiler-0.6-py3-none-any.whl size=54265 sha256=940fba6a6a6063cef9232f47a844962bf8f3f142124b4629cea6a480f4af9acbc
  Stored in directory: /root/.cache/pip/wheels/47/1c/59/b80a730f4afd2144bad854df4b167b812486c9d4c1bd4cf4c5
Successfully built skompiler
Installing collected packages: skompiler
Successfully installed skompiler-0.6
WARNING: Running pip as root will break packages and permissions. You should inst
all packages reliably by using venv: https://pip.pypa.io/warnings/venv
Note: you may need to restart the kernel to use updated packages.
```

```
In [3]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        import statsmodels.api as sm
        import statsmodels.formula.api as smf
        from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
        from sklearn.metrics import mean_squared_error, r2_score
        from sklearn.model_selection import train_test_split, cross_val_score, cross_val_f
        from sklearn.decomposition import PCA
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.preprocessing import scale
        from sklearn import model_selection
        from sklearn.metrics import roc_auc_score, roc_curve
        from sklearn import preprocessing
        from sklearn.metrics import classification_report
        from sklearn.metrics import confusion_matrix, accuracy_score
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.ensemble import RandomForestClassifier, BaseEnsemble, GradientBoostin
        from sklearn.svm import SVC, LinearSVC
        import time
        from matplotlib.colors import ListedColormap
        from xgboost import XGBRegressor
```

```
from skompiler import skompile
from lightgbm import LGBMRegressor
```

In order to see all rows and columns, we will increase max display numbers of dataframe.

```
In [4]: pd.set_option('display.max_rows', 1000)
pd.set_option('display.max_columns', 1000)
pd.set_option('display.width', 1000)
```

Support Vector Machines - Classifier(SVM) - Linear Kernel

Illustrative example:




Photo is cited by:<https://towardsdatascience.com/an-introduction-to-support-vector-regression-svr-a3ebc1672c2>

```
In [5]: df = pd.read_csv("../input/pima-indians-diabetes-database/diabetes.csv")
df.head()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreef
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
In [6]: df.shape
```

```
Out[6]: (768, 9)
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000

```
In [8]: X = df.drop("Outcome",axis=1)
y= df["Outcome"] #We will predict Outcome(diabetes)
```

Now we're going to split our dataset to train and test set. We will choose almost 20% of dataset as test size.

```
In [9]: X_train = X.iloc[:600]
X_test = X.iloc[600:]
y_train = y[:600]
y_test = y[600:]

print("X_train Shape: ",X_train.shape)
print("X_test Shape: ",X_test.shape)
print("y_train Shape: ",y_train.shape)
print("y_test Shape: ",y_test.shape)
```

```
X_train Shape: (600, 8)
X_test Shape: (168, 8)
y_train Shape: (600,)
y_test Shape: (168,)
```



```
In [10]: support_vector_classifier = SVC(kernel="linear").fit(X_train,y_train)
```

We also could use `svm.LinearSVC()` function directly.

```
In [11]: support_vector_classifier
```

```
Out[11]: SVC(kernel='linear')
```

```
In [12]: # Default C  
support_vector_classifier.C
```

```
Out[12]: 1.0
```

Prediction

```
In [13]: support_vector_classifier
```

```
Out[13]: SVC(kernel='linear')
```

Because we are doing a classification case, we will create a **confusion matrix** in order to evaluate our model.

```
In [14]: y_pred = support_vector_classifier.predict(X_test)
```

```
In [15]: cm = confusion_matrix(y_test,y_pred)
```

```
In [16]: cm
```

```
Out[16]: array([[96, 12],  
               [27, 33]])
```

- **true positive:** for correctly predicted event values.
- **false positive:** for incorrectly predicted event values.
- **true negative:** for correctly predicted no-event values.
- **false negative:** for incorrectly predicted no-event values.

		True Class	
		Positive	Negative
Predicted Class	Positive	TP	FP
	Negative	FN	TN

Photo is cited by [here](#).

```
In [17]: print("Our Accuracy is: ", (cm[0][0]+cm[1][1])/(cm[0][0]+cm[1][1]+cm[0][1]+cm[1][0]))
```

```
Our Accuracy is: 0.7678571428571429
```

```
In [18]: accuracy_score(y_test,y_pred)
```

```
Out[18]: 0.7678571428571429
```

```
In [19]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.89	0.83	108
1	0.73	0.55	0.63	60
accuracy			0.77	168
macro avg	0.76	0.72	0.73	168
weighted avg	0.76	0.77	0.76	168

Model Tuning & Validation

```
In [20]: support_vector_classifier
```

```
Out[20]: SVC(kernel='linear')
```

Now we will try to tune our model by using **K-Fold Cross Validation**.

```
In [21]: accuracies= cross_val_score(estimator=support_vector_classifier,  
                                     X=X_train,y=y_train,  
                                     cv=10)  
print("Average Accuracy: {:.2f} %".format(accuracies.mean()*100))  
print("Standart Deviation of Accuracies: {:.2f} %".format(accuracies.std()*100))
```

Average Accuracy: 77.33 %
Standart Deviation of Accuracies: 4.90 %

```
In [22]: support_vector_classifier.predict(X_test)[:10]
```

```
Out[22]: array([0, 0, 0, 1, 1, 0, 1, 0, 1, 0])
```

Now we will tune our model with GridSearch.

```
In [23]: svm_params = {"C":np.arange(1,20)}
```

```
In [24]: svm = SVC(kernel="linear")  
svm_cv = GridSearchCV(svm,svm_params,cv=8)
```

```
In [25]: start_time = time.time()  
  
svm_cv.fit(X_train,y_train)  
  
elapsed_time = time.time() - start_time  
  
print(f"Elapsed time for Support Vector Regression cross validation: "  
      f"{elapsed_time:.3f} seconds")
```

Elapsed time for Support Vector Regression cross validation: 4095.631 seconds

```
In [26]: #best score  
svm_cv.best_score_
```



```
Out[26]: 0.7716666666666667
```

```
In [27]: #best parameters  
svm_cv.best_params_
```

```
Out[27]: {'C': 2}
```

```
In [28]: svm_tuned = SVC(kernel="linear",C=2).fit(X_train,y_train)
```

```
In [29]: svm_tuned
```

```
Out[29]: SVC(C=2, kernel='linear')
```

```
In [30]: y_pred = svm_tuned.predict(X_test)
```

```
In [31]: cm = confusion_matrix(y_test,y_pred)
```

```
In [32]: cm
```

```
Out[32]: array([[96, 12],  
               [27, 33]])
```

- **true positive:** for correctly predicted event values.
- **false positive:** for incorrectly predicted event values.
- **true negative:** for correctly predicted no-event values.
- **false negative:** for incorrectly predicted no-event values.

```
In [33]: print("Our Accuracy is: ", (cm[0][0]+cm[1][1])/(cm[0][0]+cm[1][1]+cm[0][1]+cm[1][0]))
```

```
Our Accuracy is: 0.7678571428571429
```

```
In [34]: accuracy_score(y_test,y_pred)
```

```
Out[34]: 0.7678571428571429
```

```
In [35]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.89	0.83	108
1	0.73	0.55	0.63	60
accuracy			0.77	168
macro avg	0.76	0.72	0.73	168
weighted avg	0.76	0.77	0.76	168

PRACTICAL NO: 6

AIM: Adaboost Ensemble Learning

- Implement the Adaboost algorithm to create an ensemble of weak classifiers.
- Train the ensemble model on a given dataset and evaluate its performance
- Compare the results with individual weak classifiers

Requirement:

Code:

```
import pandas
from sklearn import model_selection
from sklearn.ensemble import AdaBoostClassifier
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indians-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
dataframe = pandas.read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]
seed = 7
num_trees = 30
#kfold makes trees with split number.
#kfold = model_selection.KFold(n_splits=10, random_state=seed)
#n_estimators : This is the number of trees you want to build before predictions.
#Higher number of trees give you better voting options and performance
model = AdaBoostClassifier(n_estimators=num_trees, random_state=seed)
#cross_val_score method is used to calculate the accuracy of model sliced into x, y
#cross validator cv is optional cv=kfold
results = model_selection.cross_val_score(model, X, Y)
print(results.mean())
```

Output:

```
AMD64) ] on win32
Type "help", "copyright", "credits" or "license()" for more
>>>
===== RESTART: C:\Hasan\code.py =====
0.7617774382480265
>>> |
```

PRACTICAL NO: 7

AIM: Naive Bayes' Classifier

- Implement the Naive Bayes algorithm for classification.
- Train a Naive Bayes' model using a given dataset and calculate class probabilities.
- Evaluate the accuracy of the model on test data and analyze the results.

Requirement: disease dataset

Code:

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB, CategoricalNB, GaussianNB
from sklearn.metrics import accuracy_score
import seaborn as sns
```

```
In [2]: # Load the disease dataset
df = pd.read_csv('disease.csv')
```

```
In [4]: df.head(11)
```

Out[4]:

	Sore Throat	Fever	Swollen Glands	Congestion	Headache	Diagnosis
0	Yes	Yes	Yes	Yes	Yes	Strep throat
1	No	No	No	Yes	Yes	Allergy
2	Yes	Yes	No	Yes	No	Cold
3	Yes	No	Yes	No	No	Strep throat
4	No	Yes	No	Yes	No	Cold
5	No	No	No	Yes	No	Allergy
6	No	No	Yes	No	No	Strep throat
7	Yes	No	No	Yes	Yes	Allergy
8	No	Yes	No	Yes	Yes	Cold
9	Yes	Yes	No	Yes	Yes	Cold

```
In [5]: df.tail()
```

Out[5]:

	Sore Throat	Fever	Swollen Glands	Congestion	Headache	Diagnosis
5	No	No	No	Yes	No	Allergy
6	No	No	Yes	No	No	Strep throat
7	Yes	No	No	Yes	Yes	Allergy
8	No	Yes	No	Yes	Yes	Cold
9	Yes	Yes	No	Yes	Yes	Cold

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sore Throat      10 non-null    object
1   Fever            10 non-null    object
2   Swollen Glands   10 non-null    object
3   Congestion       10 non-null    object
4   Headache         10 non-null    object
5   Diagnosis        10 non-null    object
dtypes: object(6)
memory usage: 608.0+ bytes
```

```
In [7]: #Changing the Datatypes of all the columns from object to int
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df['Sore Throat']=le.fit_transform(df['Sore Throat'])
df['Fever']=le.fit_transform(df['Fever'])
df['Swollen Glands']=le.fit_transform(df['Swollen Glands'])
df['Congestion']=le.fit_transform(df['Congestion'])
df['Headache']=le.fit_transform(df['Headache'])
df['Diagnosis']=le.fit_transform(df['Diagnosis'])
```

```
In [8]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sore Throat      10 non-null    int32
1   Fever            10 non-null    int32
2   Swollen Glands   10 non-null    int32
3   Congestion       10 non-null    int32
4   Headache         10 non-null    int32
5   Diagnosis        10 non-null    int32
dtypes: int32(6)
memory usage: 368.0 bytes
```

```
In [9]: df.head(11)
```


```
Out[9]:
```

	Sore Throat	Fever	Swollen Glands	Congestion	Headache	Diagnosis
0	1	1	1	1	1	2
1	0	0	0	1	1	0
2	1	1	0	1	0	1
3	1	0	1	0	0	2
4	0	1	0	1	0	1
5	0	0	0	1	0	0
6	0	0	1	0	0	2
7	1	0	0	1	1	0
8	0	1	0	1	1	1
9	1	1	0	1	1	1

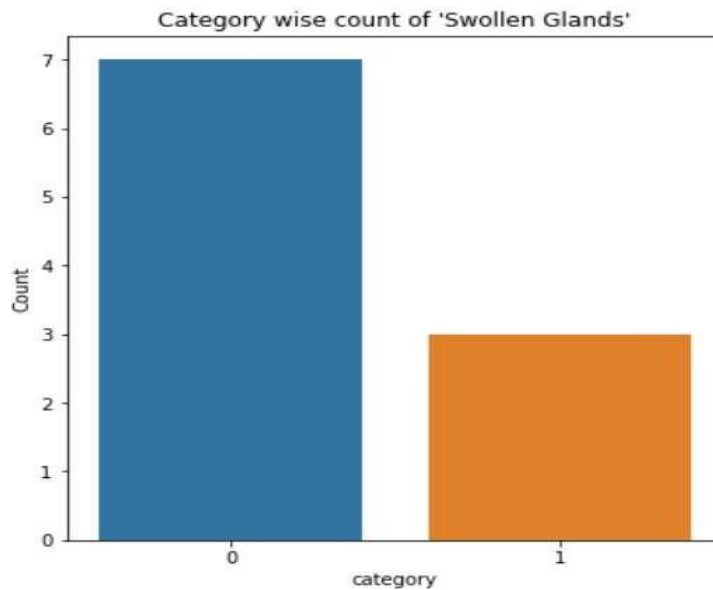
```
In [13]: #setting the dimensions of the plot
fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Sore Throat'],data=df)
plt.title("Category wise count of 'Sore Throat'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```



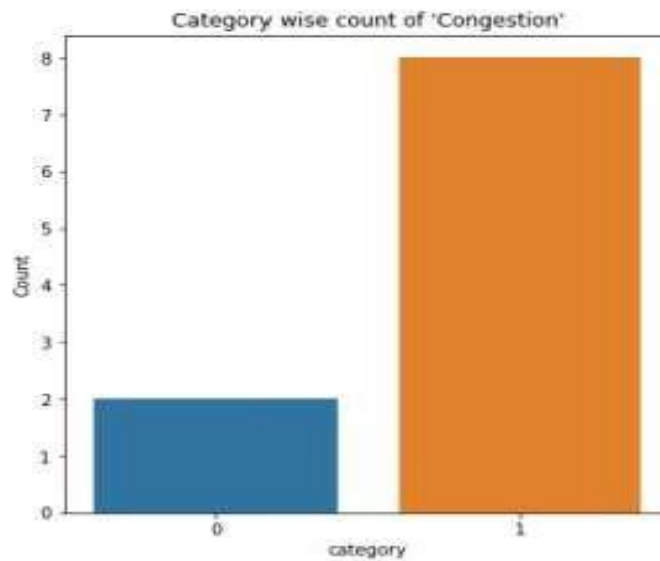
```
In [14]: fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Fever'],data=df)
plt.title("Category wise count of 'Fever'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```



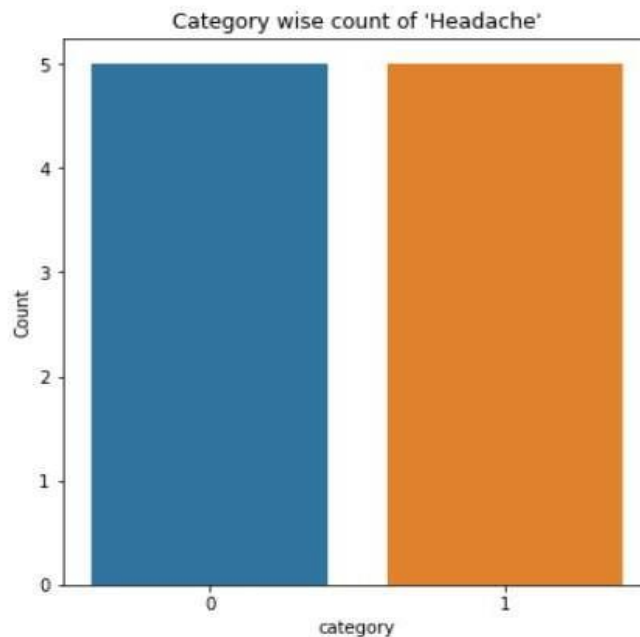
```
In [15]: fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Swollen Glands'],data=df)
plt.title("Category wise count of 'Swollen Glands'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```



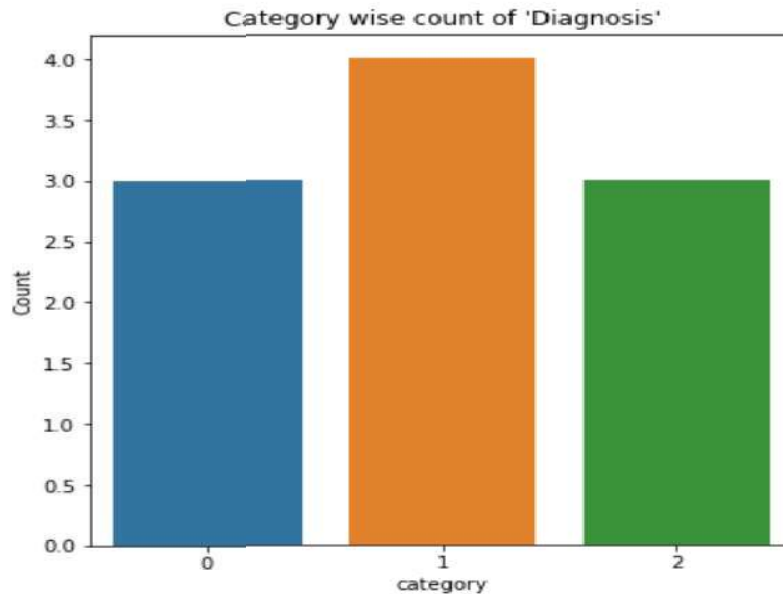
```
In [16]: fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Congestion'],data=df)
plt.title("Category wise count of 'Congestion'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```



```
In [17]: fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Headache'],data=df)
plt.title("Category wise count of 'Headache'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```




```
In [18]: fig,ax=plt.subplots(figsize=(6,6))
sns.countplot(x=df['Diagnosis'],data=df)
plt.title("Category wise count of 'Diagnosis'")
plt.xlabel("category")
plt.ylabel("Count")
plt.show()
```



```
In [19]: X=df.drop('Diagnosis',axis=1)
y=df['Diagnosis']
```

```
In [21]: #Training algorithm
classifier=MultinomialNB()
classifier.fit(X,y)
```

```
Out[21]: MultinomialNB()
```

```
In [54]: #Training algorithm
classifier=CategoricalNB()
classifier.fit(X,y)
```

```
Out[54]: CategoricalNB()
```

```
In [27]: #Training algorithm
classifier=GaussianNB()
classifier.fit(X,y)
```

```
Out[27]: GaussianNB()
```

```
In [55]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report,accuracy_score,confusion_matrix,precision_score,recall_score,f1_score
```

```
In [56]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2)
```



```
In [57]: classifier=MultinomialNB()
classifier.fit(X_train,y_train)
y_pred=classifier.predict(X_test)
print("confusion matrix\n",confusion_matrix(y_test,y_pred))
print("Accuracy:",accuracy_score(y_test,y_pred))
print("Precision:",precision_score(y_test,y_pred))
print("Recall:",recall_score(y_test,y_pred))
print("F1 score:",f1_score(y_test,y_pred))
print("Classification report:\n",classification_report(y_test,y_pred))
```

confusion matrix

```
[[1 0]
```

```
[0 1]]
```

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 score: 1.0

Classification report:\n

			precision	recall	f1-score	support
	1	1.00	1.00	1.00	1	
	2	1.00	1.00	1.00	1	
	accuracy		1.00		2	
	macro avg	1.00	1.00	1.00	2	
	weighted avg	1.00	1.00	1.00	2	

PRACTICAL NO: 8

Aim:- Implement the K-NN Algorithm for classification or regression.

Apply K-NN Algorithm on the given dataset & predict the class or value for test data.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
In [2]: df = pd.read_csv('C:/Users/RDNC/Desktop/diabetes.csv')
df.head()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

```
In [3]: df.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: df.dtypes
```

```
Out[4]: Pregnancies      int64
Glucose      int64
BloodPressure  int64
SkinThickness int64
Insulin      int64
BMI          float64
DiabetesPedigreeFunction float64
Age          int64
Outcome      int64
dtype: object
```

```
In [9]: x= df.drop('Outcome',axis=1).values
y = df['Outcome'].values
```

```
In [18]: from sklearn.model_selection import train_test_split
```

```
In [19]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.4,random_state=42, s
```

```
In [24]: from sklearn.neighbors import KNeighborsClassifier
neighbors = np.arange(1,9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):
    #Setup a knn classifier with k neighbors
    knn = KNeighborsClassifier(n_neighbors=k)
    #Fit the model
    knn.fit(X_train, y_train)
    #Compute accuracy on the training set
```

```

train_accuracy[i] = knn.score(X_train, y_train)
#Compute accuracy on the test set
test_accuracy[i] = knn.score(X_test, y_test)

```


```

In [25]: plt.title('k-NN Varying number of neighbors')
plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training accuracy')
plt.legend()
plt.xlabel('Number of neighbors')
plt.ylabel('Accuracy')
plt.show()

```

file:///C:/Users/Lenovo/Downloads/Practical no 8- K-NEAREST NEIGHBOUR.html

4/36



```

In [27]: knn = KNeighborsClassifier(n_neighbors=7)

```

```

In [30]: knn.fit(X_train,y_train)

```

```

Out[30]: KNeighborsClassifier(n_neighbors=7)

```

```

In [32]: knn.score(X_test,y_test)

```

C:\Users\RDNC\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. 'skew', 'kurtosis'), the default behavior of 'mode' typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of 'keepdims' will become False, the 'axis' over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set 'keepdims' to True or False to avoid this warning.

```

mode, _ = stats.mode(y[neigh_ind, k], axis=1)

```

```

Out[32]: 0.7142857142857143

```

```

In [36]: from sklearn.metrics import confusion_matrix

```

```

In [ ]:

```

```

In [37]: y_pred = knn.predict(X_test)

```

```
In [39]: confusion_matrix(y_test,y_pred)
```

```
Out[39]: array([[163, 43],
               [ 45, 57]], dtype=int64)
```

```
In [43]: from sklearn.metrics import classification_report
```

```
In [ ]:
```

```
In [44]: print(classification_report(y_test,y_pred))
```


	precision	recall	f1-score	support
0	0.78	0.79	0.79	206
1	0.57	0.56	0.56	102
accuracy			0.71	308
macro avg	0.68	0.68	0.68	308
weighted avg	0.71	0.71	0.71	308

```
In [46]: y_pred_proba = knn.predict_proba(X_test)[:,-1]
```

```
In [48]: from sklearn.metrics import roc_curve
```

```
In [50]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
```

```
In [52]: plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr, label='Knn')
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('Knn(n_neighbors=7) ROC curve')
plt.show()
```



```
In [54]: from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)
```

```
Out[54]: 0.7536645726251665
```

```
In [56]: from sklearn.model_selection import GridSearchCV
```

```
In [58]: param_grid = {'n_neighbors':np.arange(1,50)}
```

```
In [61]: knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(x,y)
```

```
Out[61]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
      param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
      18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
      35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49])})
```

```
In [63]: knn_cv.best_score_
```

```
Out[63]: 0.7578558696205755
```

```
In [64]: knn_cv.best_params_
```

```
Out[64]: {'n_neighbors': 14}
```